



Pipelines and AutoML with mlr3

<https://tinyurl.com/mlr3pipelines>

Department of Statistics – LMU Munich

May 28, 2020



SO YOU WANT TO DO ML IN R

- R gives you access to many machine learning methods
- ...but without a unified interface
- ...resampling and performance evaluation are cumbersome

SO YOU WANT TO DO ML IN R

- R gives you access to many machine learning methods
- ...but without a unified interface
- ...resampling and performance evaluation are cumbersome

mlr3 provides an interface to several machine learning algorithms for *training*, *predicting*, *resampling*, *tuning*, *benchmarks* and more.

SO YOU WANT TO DO ML IN R

- R gives you access to many machine learning methods
- ...but without a unified interface
- ...resampling and performance evaluation are cumbersome

mlr3 provides an interface to several machine learning algorithms for *training*, *predicting*, *resampling*, *tuning*, *benchmarks* and more.

Example:

```
data = tsk("iris")
algo = lrn("classif.ranger")
algo$train(data)
```

SO YOU WANT TO DO ML IN R

- R gives you access to many machine learning methods
- ...but without a unified interface
- ...resampling and performance evaluation are cumbersome

mlr3 provides an interface to several machine learning algorithms for *training*, *predicting*, *resampling*, *tuning*, *benchmarks* and more.

Example:

```
data = tsk("iris")
algo = lrn("classif.ranger")
algo$train(data)

algo$predict_newdata(data.frame(
  Sepal.Length = 4, Sepal.Width = 4,
  Petal.Length = 2, Petal.Width = 0.4
))

#> <PredictionClassif> for 1 observations:
#>   row_id truth response
#>     1 <NA>    setosa
```

SO YOU WANT TO DO ML IN R

- R gives you access to many machine learning methods
- ...but without a unified interface
- ...resampling and performance evaluation are cumbersome

mlr3 provides an interface to several machine learning algorithms for *training*, *predicting*, *resampling*, *tuning*, *benchmarks* and more.

Example:

```
data = tsk("iris")
algo = lrn("classif.ranger")
rr = resample(data, algo, rsmp("cv"))
rr$aggregate(msr("classif.acc"))

#> classif.acc
#>      0.96
```

SO YOU WANT TO DO ML IN R

- R gives you access to many machine learning methods
- ...but without a unified interface
- ...resampling and performance evaluation are cumbersome

mlr3 provides an interface to several machine learning algorithms for *training*, *predicting*, *resampling*, *tuning*, *benchmarks* and more.

Example:

```
design = benchmark_grid(  
  tasks = list(tsk("iris"), tsk("german_credit")),  
  learners = list(lrn("classif.ranger"), lrn("classif.rpart")),  
  resamplings = list(rsmp("cv"))  
)  
  
bmr = benchmark(design)  
bmr$aggregate(msr("classif.acc"))[,  
  .(task_id, learner_id, classif.acc)]  
  
#>      task_id    learner_id classif.acc  
#> 1:      iris classif.ranger  0.9600000  
#> 2:      iris  classif.rpart  0.9466667  
#> 3: german_credit classif.ranger  0.7720000  
#> 4: german_credit  classif.rpart  0.7310000
```

MLR3 PHILOSOPHY


- Overcome limitations of S3 with the help of **R6**
 - Truly object-oriented: data and methods live in the same object
 - Make use of inheritance
 - Reference semantics

MLR3 PHILOSOPHY

- Overcome limitations of S3 with the help of **R6**
 - Truly object-oriented: data and methods live in the same object
 - Make use of inheritance
 - Reference semantics
- Embrace **data.table**, both for arguments and internally
 - Fast operations for tabular data
 - List columns to arrange complex objects in tabular structure

MLR3 PHILOSOPHY


- Overcome limitations of S3 with the help of **R6**
 - Truly object-oriented: data and methods live in the same object
 - Make use of inheritance
 - Reference semantics
- Embrace **data.table**, both for arguments and internally
 - Fast operations for tabular data
 - List columns to arrange complex objects in tabular structure
- Be **light on dependencies**:
 - R6, data.table, Metrics, lgr, uuid, mlbench, digest
 - Plus some of our own packages (backports, checkmate, ...)



mlr3pipelines


MACHINE LEARNING WORKFLOWS

- **Preprocessing:** Feature extraction, feature selection, missing data imputation,...
- **Ensemble methods:** Model averaging, model stacking
- **mlr3:** modular model fitting
⇒ **mlr3pipelines:** modular ML workflows



MACHINE LEARNING WORKFLOWS


– what do they look like?



MACHINE LEARNING WORKFLOWS

– what do they look like?


- **Building blocks:** *what is happening? → PipeOp*



MACHINE LEARNING WORKFLOWS

– what do they look like?

- **Building blocks:** *what is happening?* → PipeOp
- **Structure:** *In what sequence is it happening?* → Graph




MACHINE LEARNING WORKFLOWS

– what do they look like?

- **Building blocks:** *what is happening?* → PipeOp

- **Structure:** In what *sequence* is it happening? → Graph


⇒ Graph: PipeOps as **nodes** with **edges** (data flow) between them



PipeOps

PIPEOP: SINGLE UNIT OF DATA OPERATION


pip = po("scale") to construct



PIPEOP: SINGLE UNIT OF DATA OPERATION

`pip$train()`: process data and create `pip$state`


\$train()




PIPEOP: SINGLE UNIT OF DATA OPERATION

`pip$predict()`: process data depending on the `pip$state`

`$train()`



`$predict()`



PIPEOP: SINGLE UNIT OF DATA OPERATION

```
po = po("scale")

trained = po$train(list(task))

trained[[1]]$head(3)

#>   Species Petal.Length Petal.Width Sepal.Length Sepal.Width
#> 1: setosa    -1.335752   -1.311052   -0.8976739   1.0156020
#> 2: setosa    -1.335752   -1.311052   -1.1392005  -0.1315388
#> 3: setosa    -1.392399   -1.311052   -1.3807271   0.3273175
```

PIPEOP: SINGLE UNIT OF DATA OPERATION

```
po = po("scale")

trained = po$train(list(task))

trained[[1]]$head(3)

#>   Species Petal.Length Petal.Width Sepal.Length Sepal.Width
#> 1: setosa    -1.335752   -1.311052   -0.8976739   1.0156020
#> 2: setosa    -1.335752   -1.311052   -1.1392005  -0.1315388
#> 3: setosa    -1.392399   -1.311052   -1.3807271   0.3273175

head(po$state, 2)

#> $center
#> Petal.Length  Petal.Width Sepal.Length  Sepal.Width
#>      3.758000    1.199333    5.843333    3.057333
#>
#> $scale
#> Petal.Length  Petal.Width Sepal.Length  Sepal.Width
#>      1.7652982   0.7622377   0.8280661   0.4358663
```

PIPEOP: SINGLE UNIT OF DATA OPERATION

```
po = po("scale")

trained = po$train(list(task))

trained[[1]]$head(3)

#>   Species Petal.Length Petal.Width Sepal.Length Sepal.Width
#> 1: setosa    -1.335752   -1.311052   -0.8976739   1.0156020
#> 2: setosa    -1.335752   -1.311052   -1.1392005  -0.1315388
#> 3: setosa    -1.392399   -1.311052   -1.3807271   0.3273175
```

```
smalltask = task$clone()$filter(1:3)
po$predict(list(smalltask))[[1]]$data()

#>   Species Petal.Length Petal.Width Sepal.Length Sepal.Width
#> 1: setosa    -1.335752   -1.311052   -0.8976739   1.0156020
#> 2: setosa    -1.335752   -1.311052   -1.1392005  -0.1315388
#> 3: setosa    -1.392399   -1.311052   -1.3807271   0.3273175
```

LIST OF PIPEOPS

Included

- Simple preprocessors (scaling, Box-Cox, Yeo-Johnson, PCA, ICA)
- NA imputation (constant, hist-sampling, model-based, dummies)
- Categorical data encoding (one-hot, treatment, impact)
- Text processing
- Feature filtering (by name, by type, statistical filters)
- Combination of data: `featureunion`
- Target column transformation (e.g. log-scaling)
- Sampling (subsampling for speed, sampling for class balance)
- Branching (simultaneous branching, alternative branching)
- Ensembling of predictions (weighted average, optimized weights)
- stacking (see later slides)


Planned

- Time series and spatio-temporal data
- Multi-output and ordinal targets

Graph Operations


GRAPH OPERATIONS

`%>>%` concatenates Graphs and PipeOps



GRAPH OPERATIONS


`gunion()` unites Graphs and PipeOps



Linear Pipelines


LINEAR PREPROCESSING

```
graph_pp = po("scale") %>>%
  po("encode") %>>%
  po("imputemedian") %>>%
  lrn("classif.rpart")
```




LINEAR PREPROCESSING

- `train()`ing: Data propagates and creates \$states




LINEAR PREPROCESSING

- `train()`ing: Data propagates and creates \$states




LINEAR PREPROCESSING

- `train()`ing: Data propagates and creates \$states




LINEAR PREPROCESSING

- `train()`ing: Data propagates and creates \$states




LINEAR PREPROCESSING

- `train()`ing: Data propagates and creates \$states



LINEAR PREPROCESSING

- `train()`ing: Data propagates and creates `$states`
- `predict()`ition: Data propagates, uses `$states`



LINEAR PREPROCESSING

```
scale %>>% encode %>>% impute %>>% rpart
```

- Setting / retrieving parameters: \$param_set

```
graph_pp$pipeops$scale$param_set$values$center = FALSE
```

LINEAR PREPROCESSING

```
scale %>>% encode %>>% impute %>>% rpart
```

- Setting / retrieving parameters: `$param_set`

```
graph_pp$pipeops$scale$param_set$values$center = FALSE
```

- Retrieving state: `$state` of individual PipeOps (*after \$train()*)

```
graph_pp$pipeops$scale$state$scale
#> Petal.Length  Petal.Width Sepal.Length  Sepal.Width
#>     4.163367     1.424451     5.921098     3.098387
```

LINEAR PREPROCESSING

```
scale %>>% encode %>>% impute %>>% rpart
```

- Setting / retrieving parameters: `$param_set`

```
graph_pp$pipeops$scale$param_set$values$center = FALSE
```

- Retrieving state: `$state` of individual PipeOps (*after \$train()*)


```
graph_pp$pipeops$scale$state$scale  
#> Petal.Length  Petal.Width Sepal.Length  Sepal.Width  
#>     4.163367    1.424451    5.921098    3.098387
```

- Retrieving intermediate results: `$.result` (set debug option before)

```
graph_pp$pipeops$scale$.result[[1]]$head(3)  
#>   Species Petal.Length Petal.Width Sepal.Length Sepal.Width  
#> 1:  setosa    0.3362663    0.140405    0.8613268    1.1296201  
#> 2:  setosa    0.3362663    0.140405    0.8275493    0.9682458  
#> 3:  setosa    0.3122473    0.140405    0.7937718    1.0327956
```


Nonlinear Pipelines

PIPEOPS WITH MULTIPLE INPUTS / OUTPUTS




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")
```




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")  
  
graph_bag = ppl("grepligate", single_path, n = 3) %>>%  
  po("classifavg")
```




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")  
  
graph_bag = ppl("gre replicate", single_path, n = 3) %>>%  
  po("classifavg")
```




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")  
  
graph_bag = ppl("grepligate", single_path, n = 3) %>>%  
  po("classifavg")
```




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")  
  
graph_bag = ppl("grepligate", single_path, n = 3) %>>%  
  po("classifavg")
```




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")  
  
graph_bag = ppl("grepligate", single_path, n = 3) %>>%  
  po("classifavg")
```




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")  
  
graph_bag = ppl("grepligate", single_path, n = 3) %>>%  
  po("classifavg")
```




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")  
  
graph_bag = ppl("grepligate", single_path, n = 3) %>>%  
  po("classifavg")
```




ENSEMBLE METHOD: BAGGING

```
single_path = po("subsample") %>>% lrn("classif.rpart")  
  
graph_bag = ppl("grepligate", single_path, n = 3) %>>%  
  po("classifavg")
```



ENSEMBLE METHOD: STACKING


```
graph_stack = gunion(list(  
    po("learner_cv", learner = lrn("regr.lm")),  
    po("learner_cv", learner = lrn("regr.svm")),  
    po("nop")))%>>%  
po("featureunion")%>>%  
lrn("regr.ranger")
```



BRANCHING

```
graph_branch = po("branch", c("pca", "ica")) %>>%
  gunion(list(po("pca"), po("ica"))) %>>%
  po("unbranch", c("pca", "ica")) %>>%
  lrn("classif.kknn")
```


Execute only one of several alternative paths



BRANCHING

```
graph_branch = po("branch", c("pca", "ica")) %>>%
  gunion(list(po("pca"), po("ica"))) %>>%
  po("unbranch", c("pca", "ica")) %>>%
  lrn("classif.kknn")
```


```
> graph_branch$pipeops$branch$  
  param_set$values$selection = "pca"
```



BRANCHING

```
graph_branch = po("branch", c("pca", "ica")) %>>%
  gunion(list(po("pca"), po("ica"))) %>>%
  po("unbranch", c("pca", "ica")) %>>%
  lrn("classif.kknn")
```

```
> graph_branch$pipeops$branch$  
  param_set$values$selection = "ica"
```




BRANCHING

Alternative:

```
graph_branch = ppl("branch",
  list(pca = po("pca"), ica = po("ica"))) %>>%
  lrn("classif.kknn")
```

```
> graph_branch$pipeops$branch$  
  param_set$values$selection = "ica"
```



Targeting Columns

RESTRICT PIPEOPS TO COLS WITH SELECTORS

Suppose we only want PCA on some columns of our data:

```
task$data(1:9)

#>      Species Petal.Length Petal.Width Sepal.Length Sepal.Width
#> 1:  setosa       1.4        0.2       5.1       3.5
#> 2:  setosa       1.4        0.2       4.9       3.0
#> 3:  setosa       1.3        0.2       4.7       3.2
#> 4:  setosa       1.5        0.2       4.6       3.1
#> 5:  setosa       1.4        0.2       5.0       3.6
#> 6:  setosa       1.7        0.4       5.4       3.9
#> 7:  setosa       1.4        0.3       4.6       3.4
#> 8:  setosa       1.5        0.2       5.0       3.4
#> 9:  setosa       1.4        0.2       4.4       2.9
```

RESTRICT PIPEOPS TO COLUMNS WITH SELECTORS

Option 1: PipeOps affect_columns parameter

```
my_pca = po("pca", affect_columns = selector_grep("^Sepal"))


result = my_pca$train(list(task))

result[[1]]$data(1:3)

#>   Species          PC1          PC2 Petal.Length Petal.Width
#> 1:  setosa -0.7781478  0.37813255      1.4        0.2
#> 2:  setosa -0.9350903 -0.13700728      1.4        0.2
#> 3:  setosa -1.1513076  0.04533873      1.3        0.2
```

RESTRICT PIPEOPS TO COLS WITH SELECTORS

Option 2: Use `po("select")`



```
sel1 = selector_grep("^Sepal")
sel2 = selector_invert(sel1)

my_pca = gunion(list(
  po("select", selector = sel1) %>>% po("pca"),
  po("select", selector = sel2, id = "select2"))
) %>>% po("featureunion")

my_pca$train(task)[[1]]
```

Having trouble remembering these?

“Pipelines” Dictionary & Short Form

“PIPELINES” DICTIONARY & SHORT FORM

Many frequently used *patterns* for pipelines

- Making Learners robust to bad data (imputation + feature encoding + ...)
- Bagging
- Branching

“PIPELINES” DICTIONARY & SHORT FORM

Many frequently used *patterns* for pipelines

- Making Learners robust to bad data (imputation + feature encoding + ...)
- Bagging
- Branching

Collection of these is in `mlr3pipelines`

`ppl()` accesses the `mlr_graphs` “Dictionary” of pre-constructed partial Graphs.

```
head(as.data.table(mlr_graphs), 5)

#>           key
#> 1:    bagging
#> 2:    branch
#> 3:  greplicate
#> 4:  robustify
#> 5: targettrafo
```


“PIPELINES” DICTIONARY & SHORT FORM

Many frequently used *patterns* for pipelines

- Making Learners robust to bad data (imputation + feature encoding + ...)
- Bagging
- Branching

Collection of these is in `mlr3pipelines`

```
gr = ppl("robustify")
plot(gr)
```




"PIPELINES" DICTIONARY & SHORT FORM

Many frequently used *patterns* for pipelines

- Making Learners robust to bad data (imputation + feature encoding + ...)
- Bagging
- Branching

Collection of these is in `mlr3pipelines`

```
gr = ppl("branch", list(po("pca"), po("nop")))
plot(gr)
```



AutoML with ‘mlr3pipelines’

AUTOML <3 PIPELINES

- AutoML: Automatic Machine Learning

AUTOML <3 PIPELINES

- AutoML: Automatic Machine Learning
- Let the algorithm make decisions about
 - ➊ *what learner to use,*
 - ➋ *what preprocessing to use, and*
 - ➌ *what hyperparameters to use.*

AUTOML <3 PIPELINES

- AutoML: Automatic Machine Learning
- Let the algorithm make decisions about
 - ➊ *what learner to use,*
 - ➋ *what preprocessing to use, and*
 - ➌ *what hyperparameters to use.*
- (1) and (2) are decisions about *graph structure* in `mlr3pipelines`


AUTOML <3 PIPELINES

- AutoML: Automatic Machine Learning
 - Let the algorithm make decisions about
 - ➊ *what learner to use,*
 - ➋ *what preprocessing to use, and*
 - ➌ *what hyperparameters to use.*
 - (1) and (2) are decisions about *graph structure* in `mlr3pipelines`
- ⇒ The problem reduces to **pipelines + parameter tuning**

AUTOML WITH MLR3PIPELINES


AutoML in a Nutshell

- Preprocessing steps
- ML Algorithms
- Tuner




GRAPHLEARNER

- Graph as a Learner
- All benefits of mlr3: **resampling, tuning, nested resampling, ...**



```
graph_pp = po("scale") %>>% po("encode") %>>%  
  po("imputemedian") %>>% lrn("classif.rpart")  
glrn = GraphLearner$new(graph_pp)  
glrn$train(task)  
glrn$predict(task)  
resample(task, glrn, rsmp("cv", folds = 3))
```

TUNING




PIPELINES TUNING

- Works **exactly** as in basic `mlr3` / `mlr3tuning`
- PipeOps have *hyperparameters* (using `paradox` pkg)
- Graphs have hyperparameters of all components *combined*
- ⇒ Joint **tuning** and nested CV of complete graph

```
p1 = ppl("branch", list(  
  "pca" = po("pca"),  
  "nothing" = po("nop")))  
  
p2 = flt("anova")  
  
p3 = ppl("branch", list(  
  "svm" = lrn("classif.svm", id = "svm", kernel = "radial"),  
  "xgb" = lrn("classif.xgboost", id = "xgb"),  
  "rf" = lrn("classif.ranger", id = "rf")), prefix_branchops = "lrn_")  
gr = p1 %>>% p2 %>>% p3  
glrn = GraphLearner$new(gr)
```

PIPELINES TUNING



PIPELINES TUNING

```
ps = ParamSet$new(list(
  ParamFct$new("branch.selection", levels = c("pca", "nothing")),
  ParamDbl$new("anova.filter.frac", lower = 0.1, upper = 1),
  ParamFct$new("lrn_branch.selection", levels = c("svm", "xgb", "rf")),
  ParamInt$new("rf.mtry", lower = 1L, upper = 20L),
  ParamInt$new("xgb.nrounds", lower = 1, upper = 500),
  ParamDbl$new("svm.cost", lower = -12, upper = 4),
  ParamDbl$new("svm.gamma", lower = -12, upper = -1)))
ps$add_dep("rf.mtry", "lrn_branch.selection", CondEqual$new("rf"))
ps$add_dep("xgb.nrounds", "lrn_branch.selection", CondEqual$new("xgb"))
ps$add_dep("svm.cost", "lrn_branch.selection", CondEqual$new("svm"))
ps$add_dep("svm.gamma", "lrn_branch.selection", CondEqual$new("svm"))
ps$trafo = function(x, param_set) {
  if (x$lrn_branch.selection == "svm")
    x$svm.cost = 2^x$svm.cost; x$svm.gamma = 2^x$svm.gamma
  return(x)
}
inst = TuningInstance$new(tsk("sonar"), glrn, rsmp("cv", iters=3),
  msr("classif.ce"), ps, term("evals", n_evals = 10))
tnr("random_search")$tune(inst)
```

mlr3(pipelines) Resources

MLR3(PIPELINES) RESOURCES

mlr3 book

The screenshot shows the mlr3 book's Pipelines chapter. The left sidebar contains a navigation menu with sections like Introduction and Overview, Basics, Model Optimization, Pipelines, and Technical. The main content area is titled "4 Pipelines". It discusses mlr3pipelines as a dataflow programming toolkit and provides an example of a pipeline graph:

```
graph LR; Scaling --> FactorEncoding; FactorEncoding --> MedianImputation; MedianImputation --> Learner
```

Below the graph, it says: "Single computational steps can be represented as so-called PipeOps, which can then be connected with directed edges in a Graph. The scope of `mlr3pipelines` is still growing. Currently supported features are:

<https://mlr3book.mlr-org.com/>

mlr3 Use Case “Gallery”

The screenshot shows the mlr3 gallery. It lists several use cases:

- mlr3 and OpenML - Moneyball use case
- A pipeline for the titanic data set - Advanced
- Tuning a stacked learner

Each use case has a brief description and a corresponding visualization or diagram.

<https://mlr3gallery.mlr-org.com/>

“cheat sheets”

The screenshot shows three cheat sheets from the mlr3 website:

- Machine learning with mlr3 :: CHEAT SHEET**: A general overview of machine learning with mlr3.
- Hyperparameter Tuning with mlr3tuning :: CHEAT SHEET**: A guide to hyperparameter tuning.
- Dataflow programming with mlr3pipelines :: CHEAT SHEET**: A guide to dataflow programming with mlr3pipelines.

<https://cheatsheets.mlr-org.com/>

OUTLOOK

What is to come?

- `mlr3pipelines`: caching, parallelization
- Better **tuners**: Bayesian Optimization, Hyperband
- Survival and Forecasting (via `mlr3proba`, `mlr3forecast`)
- Deep Learning (via `mlr3keras`)

Thanks! Please ask questions!